

Shiny Sherly Katuru

7007443414

Neural Networks and Deep Learning

Assignment-7

In class programming:

1. Follow the instruction below and then report how the performance changed.(apply all at once)
 - Convolutional input layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Flatten layer.
 - Dropout layer at 20%.
 - Fully connected layer with 1024 units and a rectifier activation function.
 - Dropout layer at 20%.
 - Fully connected layer with 512 units and a rectifier activation function.
 - Dropout layer at 20%.
 - Fully connected output layer with 10 units and a Softmax activation function Did the performance change?

Home Page - Select or create a notebook | ICP7_NN - Jupyter Notebook | +

localhost:8888/notebooks/ICP7_NN.ipynb

jupyter ICP7_NN Last Checkpoint: 3 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

```
In [13]: import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.constraints import maxnorm
from keras.optimizers import SGD
from keras.layers.convolutional import Conv2D, MaxPooling2D
from keras.utils import np_utils

In [3]: np.random.seed(7)

In [4]: (X_train, y_train), (X_test, y_test) = cifar10.load_data()

In [5]: X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

In [6]: y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
num_classes = y_test.shape[1]

In [7]: model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2), padding='same'))
model.add(Flatten())
model.add(Dense(512, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
```

52°F Cloudy 11:58 PM 3/22/2023

Home Page - Select or create a notebook | ICP7_NN - Jupyter Notebook | +

localhost:8888/notebooks/ICP7_NN.ipynb

jupyter ICP7_NN Last Checkpoint: 3 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

```
In [7]: model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2), padding='same'))
model.add(Flatten())
model.add(Dense(512, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

In [8]: sgd = SGD(learning_rate=0.01, momentum=0.9, decay=1e-6)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

Model: "sequential"

```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
dropout (Dropout)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 512)	4194816

52°F Cloudy 11:58 PM 3/22/2023

Home Page - Select or create a notebook | ICP7_NN - Jupyter Notebook

localhost:8888/notebooks/ICP7_NN.ipynb

jupyter ICP7_NN Last Checkpoint: 3 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

```
Total params: 4,210,090
Trainable params: 4,210,090
Non-trainable params: 0

None

In [9]: epochs = 5
        batch_size = 32
        model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=batch_size)

Epoch 1/5
1563/1563 [=====] - 178s 113ms/step - loss: 1.6913 - accuracy: 0.3858 - val_loss: 1.4168 - val_accuracy: 0.5036
Epoch 2/5
1563/1563 [=====] - 174s 111ms/step - loss: 1.3796 - accuracy: 0.5033 - val_loss: 1.2318 - val_accuracy: 0.5636
Epoch 3/5
1563/1563 [=====] - 173s 111ms/step - loss: 1.2227 - accuracy: 0.5623 - val_loss: 1.1281 - val_accuracy: 0.5998
Epoch 4/5
1563/1563 [=====] - 181s 116ms/step - loss: 1.0890 - accuracy: 0.6138 - val_loss: 1.0736 - val_accuracy: 0.6213
Epoch 5/5
1563/1563 [=====] - 181s 116ms/step - loss: 0.9815 - accuracy: 0.6542 - val_loss: 1.0174 - val_accuracy: 0.6504

Out[9]: <keras.callbacks.History at 0x29d802dc220>

In [10]: scores = model.evaluate(X_test, y_test, verbose=0)
         print("Accuracy: %.2f%%" % (scores[1]*100))

Accuracy: 65.04%
```

52°F Cloudy Search 11:58 PM 3/22/2023

Home Page - Select or create a notebook | ICP7_NN - Jupyter Notebook | +

localhost:8888/notebooks/ICP7_NN.ipynb

jupyter ICP7_NN Last Checkpoint: 3 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

```
In [10]: scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))

Accuracy: 65.04%

In [11]: import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers.convolutional import Conv2D, MaxPooling2D
from keras.constraints import maxnorm
from keras.utils import np_utils
from keras.optimizers import SGD

# Fix random seed for reproducibility
np.random.seed(7)

# Load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

# Normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# One hot encode outputs
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
num_classes = y_test.shape[1]

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
```

52°F Cloudy 11:58 PM 3/22/2023

Home Page - Select or create a notebook | ICP7_NN - Jupyter Notebook | +

localhost:8888/notebooks/ICP7_NN.ipynb

jupyter ICP7_NN Last Checkpoint: 3 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

```
num_classes = y_test.shape[1]

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 5
learning_rate = 0.01
decay_rate = learning_rate / epochs
sgd = SGD(lr=learning_rate, momentum=0.9, decay=decay_rate, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)
```

52°F Cloudy 11:58 PM 3/22/2023

Home Page - Select or create a notebook | ICP7_NN - Jupyter Notebook | localhost:8888/notebooks/ICP7_NN.ipynb

jupyter ICP7_NN Last Checkpoint: 3 hours ago (autosaved) | Python 3 (ipykernel)

```
# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)

# Evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 32, 32, 32)	896
dropout_2 (Dropout)	(None, 32, 32, 32)	0
conv2d_3 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_4 (Conv2D)	(None, 16, 16, 64)	18496
dropout_3 (Dropout)	(None, 16, 16, 64)	0
conv2d_5 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_6 (Conv2D)	(None, 8, 8, 128)	73856
dropout_4 (Dropout)	(None, 8, 8, 128)	0

52°F Cloudy | 11:58 PM 3/22/2023

Home Page - Select or create a notebook | ICP7_NN - Jupyter Notebook | localhost:8888/notebooks/ICP7_NN.ipynb

jupyter ICP7_NN Last Checkpoint: 3 hours ago (autosaved) | Python 3 (ipykernel)

```
dense_4 (Dense) (None, 10) 5130

Total params: 2,915,114
Trainable params: 2,915,114
Non-trainable params: 0
```

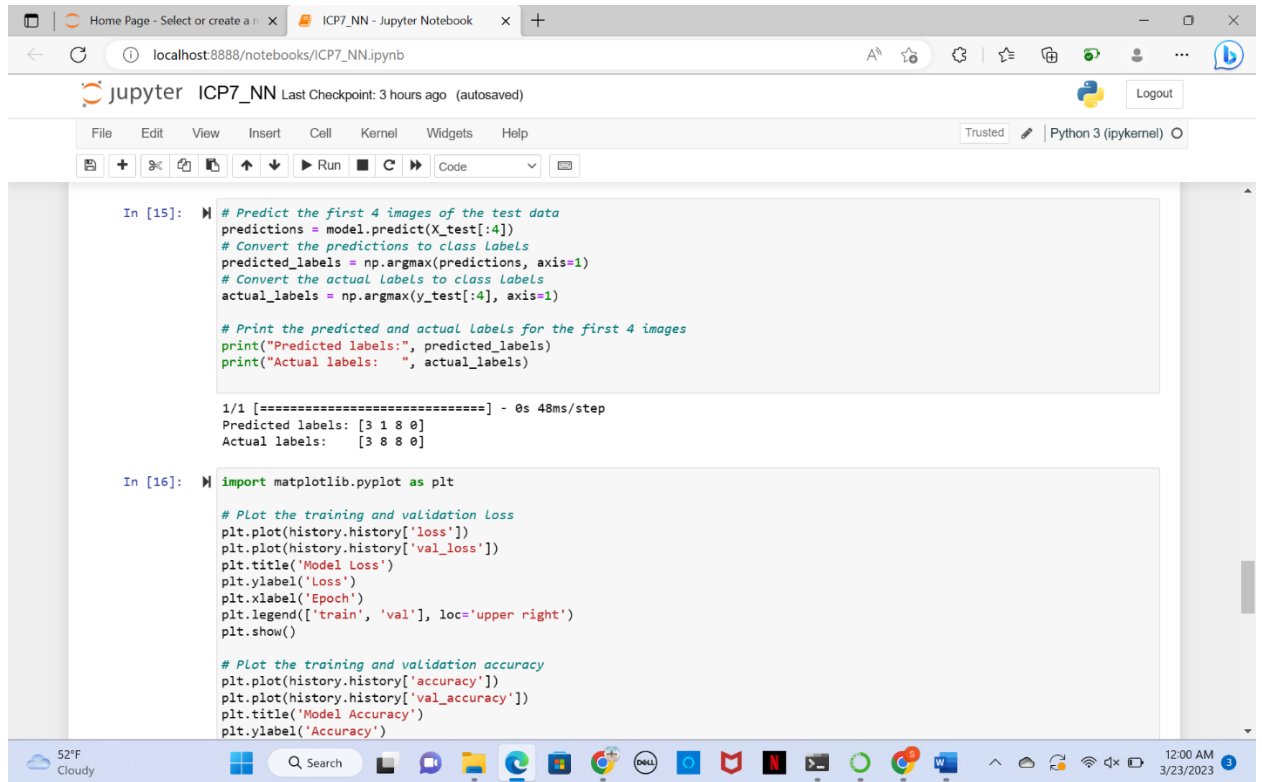
C:\Users\shiny\anaconda3\lib\site-packages\keras\optimizers\optimizer_v2\gradient_descent.py:114: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
super().__init__(name, **kwargs)

```
None
Epoch 1/5
1563/1563 [=====] - 349s 222ms/step - loss: 1.8774 - accuracy: 0.3035 - val_loss: 1.5661 - val_accuracy: 0.4307
Epoch 2/5
1563/1563 [=====] - 1512s 968ms/step - loss: 1.5070 - accuracy: 0.4497 - val_loss: 1.4609 - val_accuracy: 0.4680
Epoch 3/5
1563/1563 [=====] - 342s 219ms/step - loss: 1.3795 - accuracy: 0.4988 - val_loss: 1.3096 - val_accuracy: 0.5230
Epoch 4/5
1563/1563 [=====] - 343s 219ms/step - loss: 1.2962 - accuracy: 0.5332 - val_loss: 1.2196 - val_accuracy: 0.5589
Epoch 5/5
1563/1563 [=====] - 340s 217ms/step - loss: 1.2410 - accuracy: 0.5519 - val_loss: 1.2445 - val_accuracy: 0.5544
Accuracy: 55.44%
```

In [15]: # Predict the first 4 images of the test data
predictions = model.predict(X_test[:4])
Convert the predictions to class labels

52°F Cloudy | 12:00 AM 3/23/2023

2. Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4 images to check whether or not the model has predicted correctly.



The screenshot shows a Jupyter Notebook interface with two code cells. The first cell, labeled 'In [15]:', contains Python code to predict the first 4 images of the test data, convert predictions to class labels, and compare them with actual labels. The output shows 'Predicted labels: [3 1 8 0]' and 'Actual labels: [3 8 8 0]'. The second cell, labeled 'In [16]:', contains code to plot the training and validation loss and accuracy using matplotlib. The output shows the plots are displayed.

```
In [15]: # Predict the first 4 images of the test data
predictions = model.predict(X_test[:4])
# Convert the predictions to class Labels
predicted_labels = np.argmax(predictions, axis=1)
# Convert the actual Labels to class Labels
actual_labels = np.argmax(y_test[:4], axis=1)

# Print the predicted and actual Labels for the first 4 images
print("Predicted labels:", predicted_labels)
print("Actual labels: ", actual_labels)

1/1 [=====] - 0s 48ms/step
Predicted labels: [3 1 8 0]
Actual labels:    [3 8 8 0]

In [16]: import matplotlib.pyplot as plt

# Plot the training and validation Loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'val'], loc='upper right')
plt.show()

# Plot the training and validation accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
```

3. Visualize Loss and Accuracy using the history object

